

DA343A

Objektorienterad programutveckling, trådar och datakommunikation

Föreläsning 9 – Synchronization

Ben Blamey

Föreläsning 9 : Synchronization

- Revision
 - Callbacks
 - Communication between threads
- Why do we need synchronization? – Method Invocation Ordering
- Synchronization
 - The synchronized keyword
 - wait() / notifyAll()
 - Synchronized blocks
- Thread termination
- Queues & Buffers
- Multi-Threaded Producer/Consumer Example

Revision: Callbacks


- A callback is a method that is registered with another object. The callback method can be invoked when:
 - A task is completed (or progress, error),
 - Something changes in the object.
- Callback can be invoked: immediately (synchronously) or later (asynchronously).
- We can use callbacks on a single thread, or with multiple threads.



Revision: Communication Between Threads

- **No communication.** “fire and forget”
- **Polling** (not recommended)
- **Callbacks** (can use arguments for extra information, ideally immutable types) (F8)
- **Shared State** – this requires some form of **synchronization**. Why?

Event Ordering Matters!



4 Events:

- (A) Train 1 enters single track section.
- (B) Train 1 leaves single track section.
- (C) Train 2 enters single track section.
- (D) Train 2 leaves single track section.

(A)	⇒	(B)	⇒	(C)	⇒	(D)	OK
(C)	⇒	(D)	⇒	(A)	⇒	(B)	OK
(A)	⇒	(C)	⇒	()	⇒	()	Collision!
(C)	⇒	(A)	⇒	()	⇒	()	Collision!

```
class SharedState {
    int x = 0;

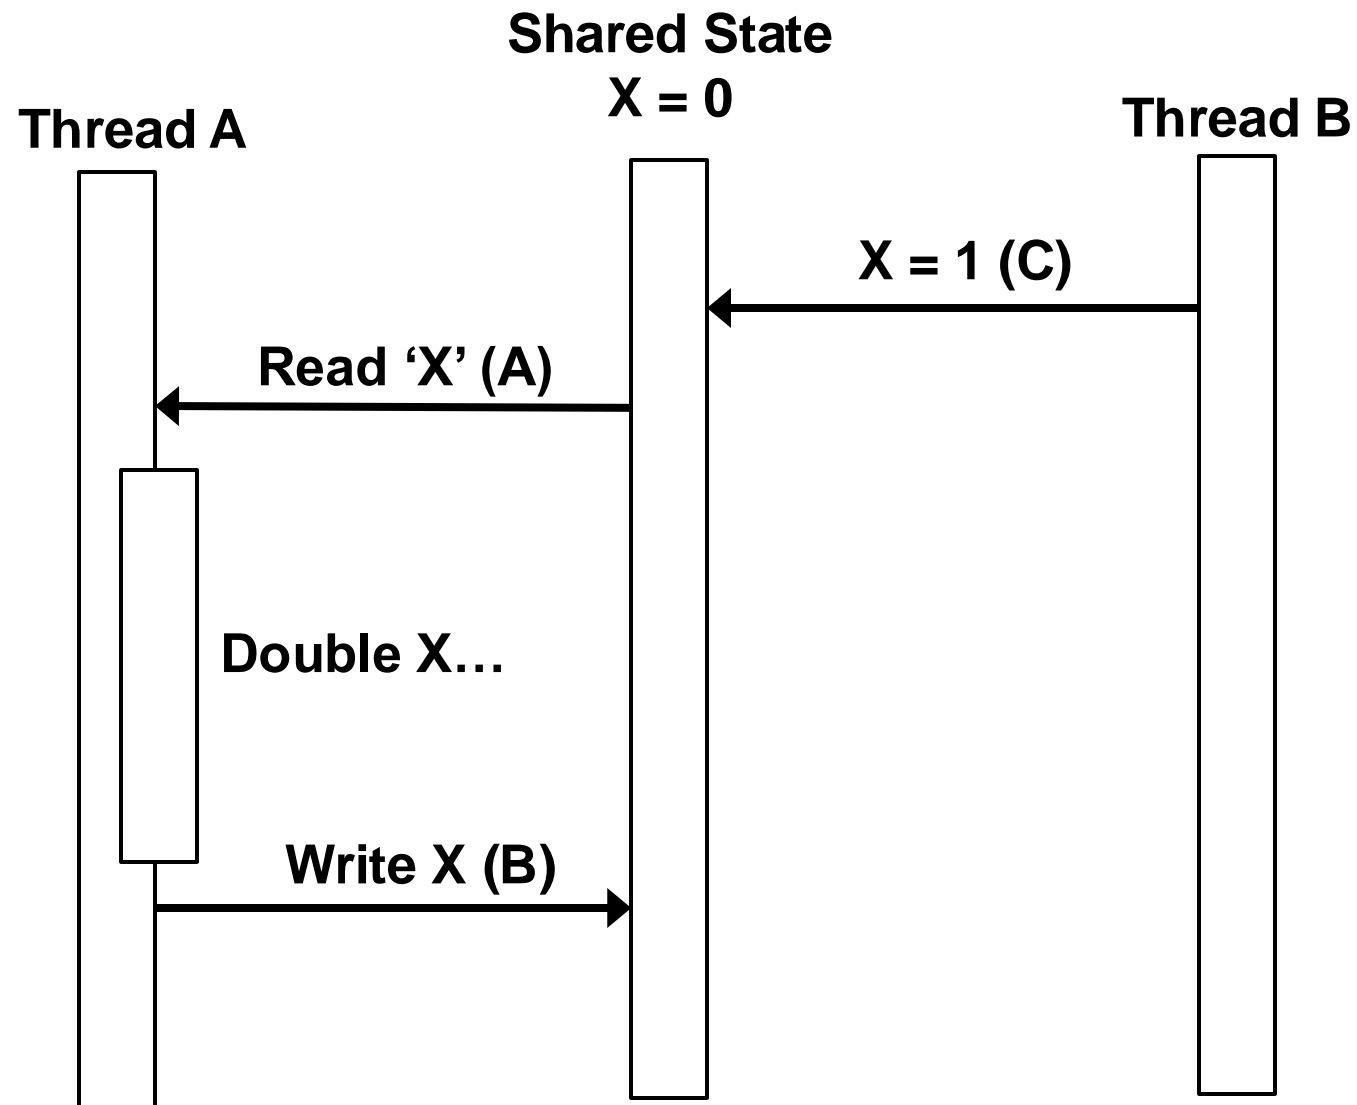
    public void doubleX() {
        int x = this.x;
        this.x = x * 2;
    }

    public void setX1() {
        this.x = 1;
    }
}

private void run() throws InterruptedException {
    SharedState someSharedState = new SharedState();
    Thread threadA = new Thread(() -> {
        someSharedState.doubleX();
    });
    Thread threadB = new Thread(() -> {
        someSharedState.setX1();
    });

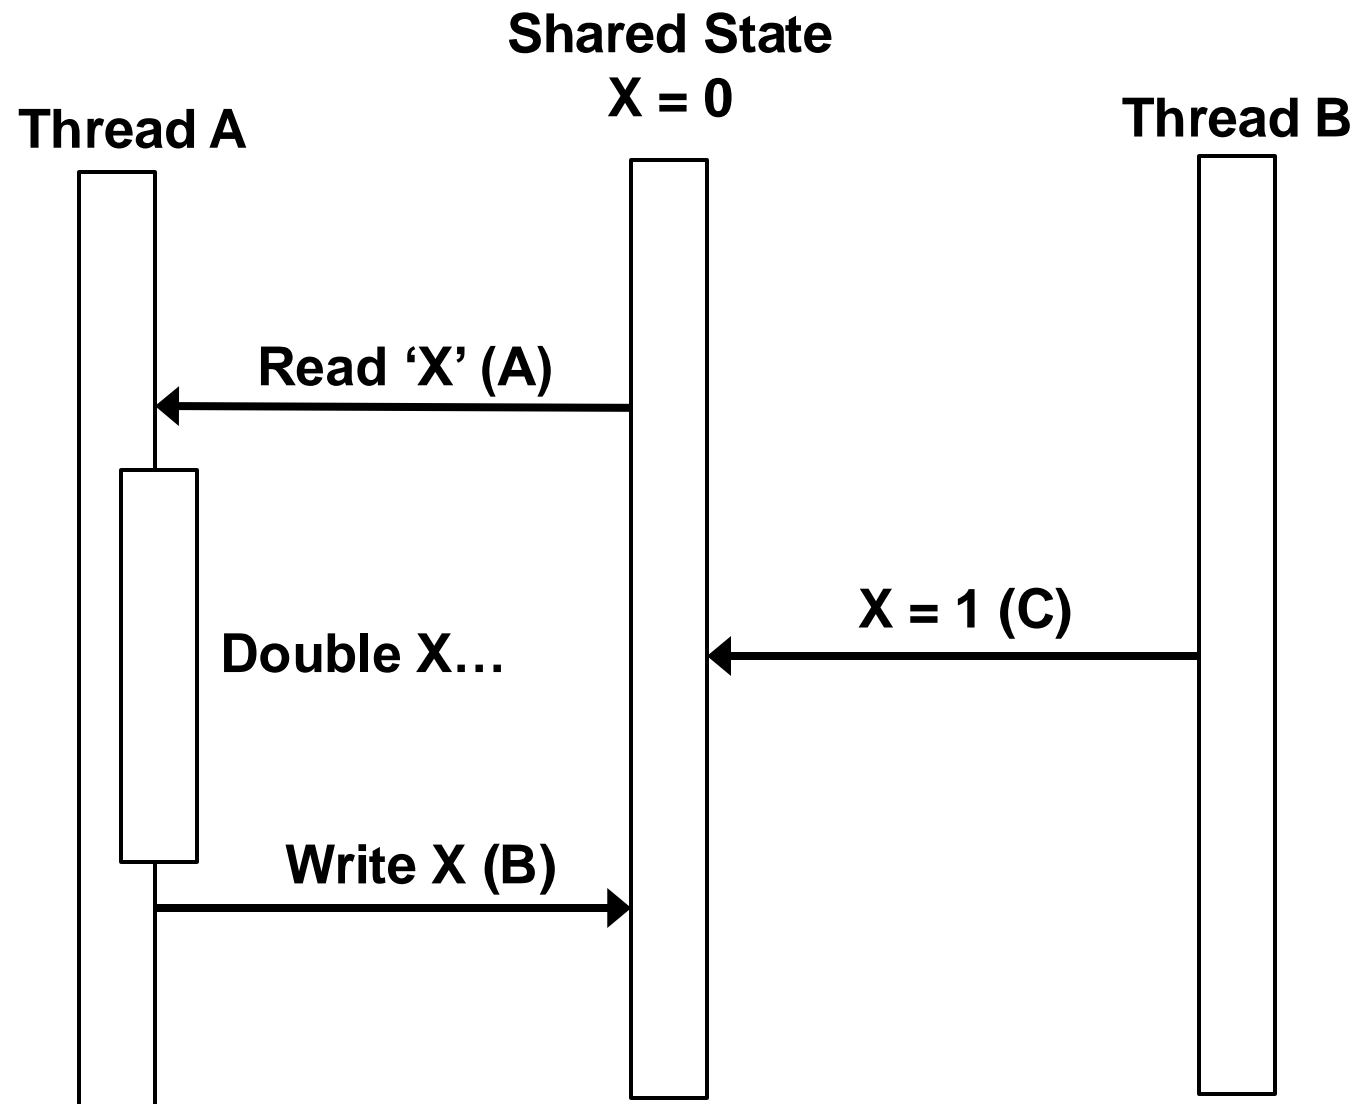
    threadA.start(); threadB.start();
    threadA.join(); threadB.join();
    System.out.print(String.format("X = %d . ", someSharedState.x));
}

public static void main(String args[]) throws InterruptedException {
    for (int i = 0; i < 100; i++) {
        new BrokenExample().run();
    }
}
```



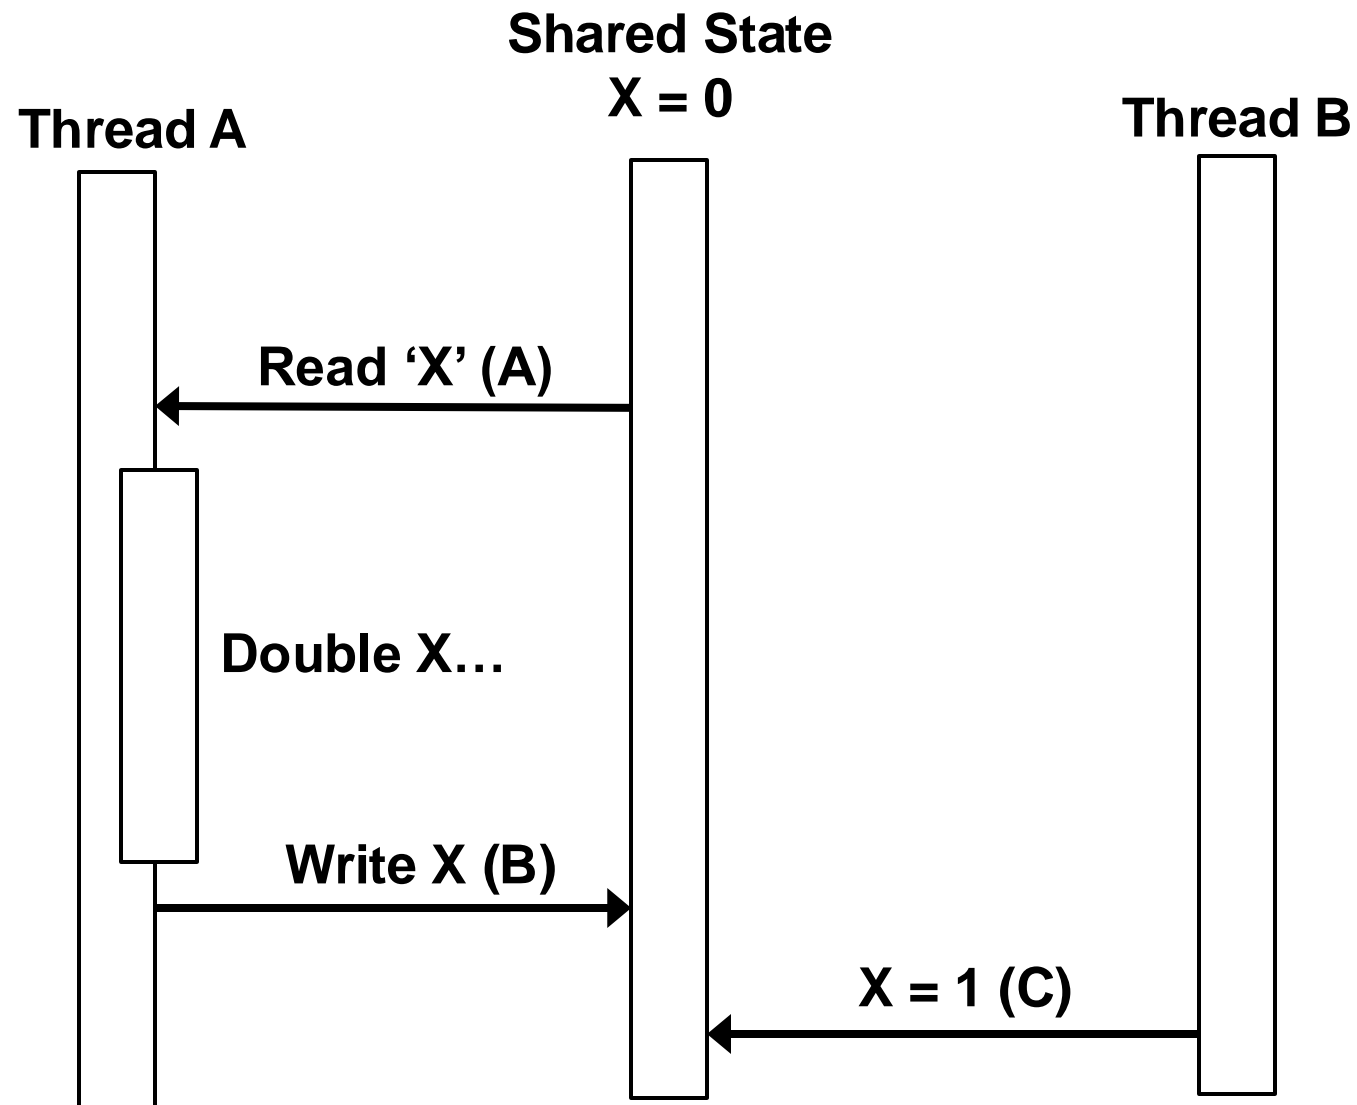
In which order does it happen?
What is the final value of X in each case?

$C \Rightarrow A \Rightarrow B$
 $X = 2$



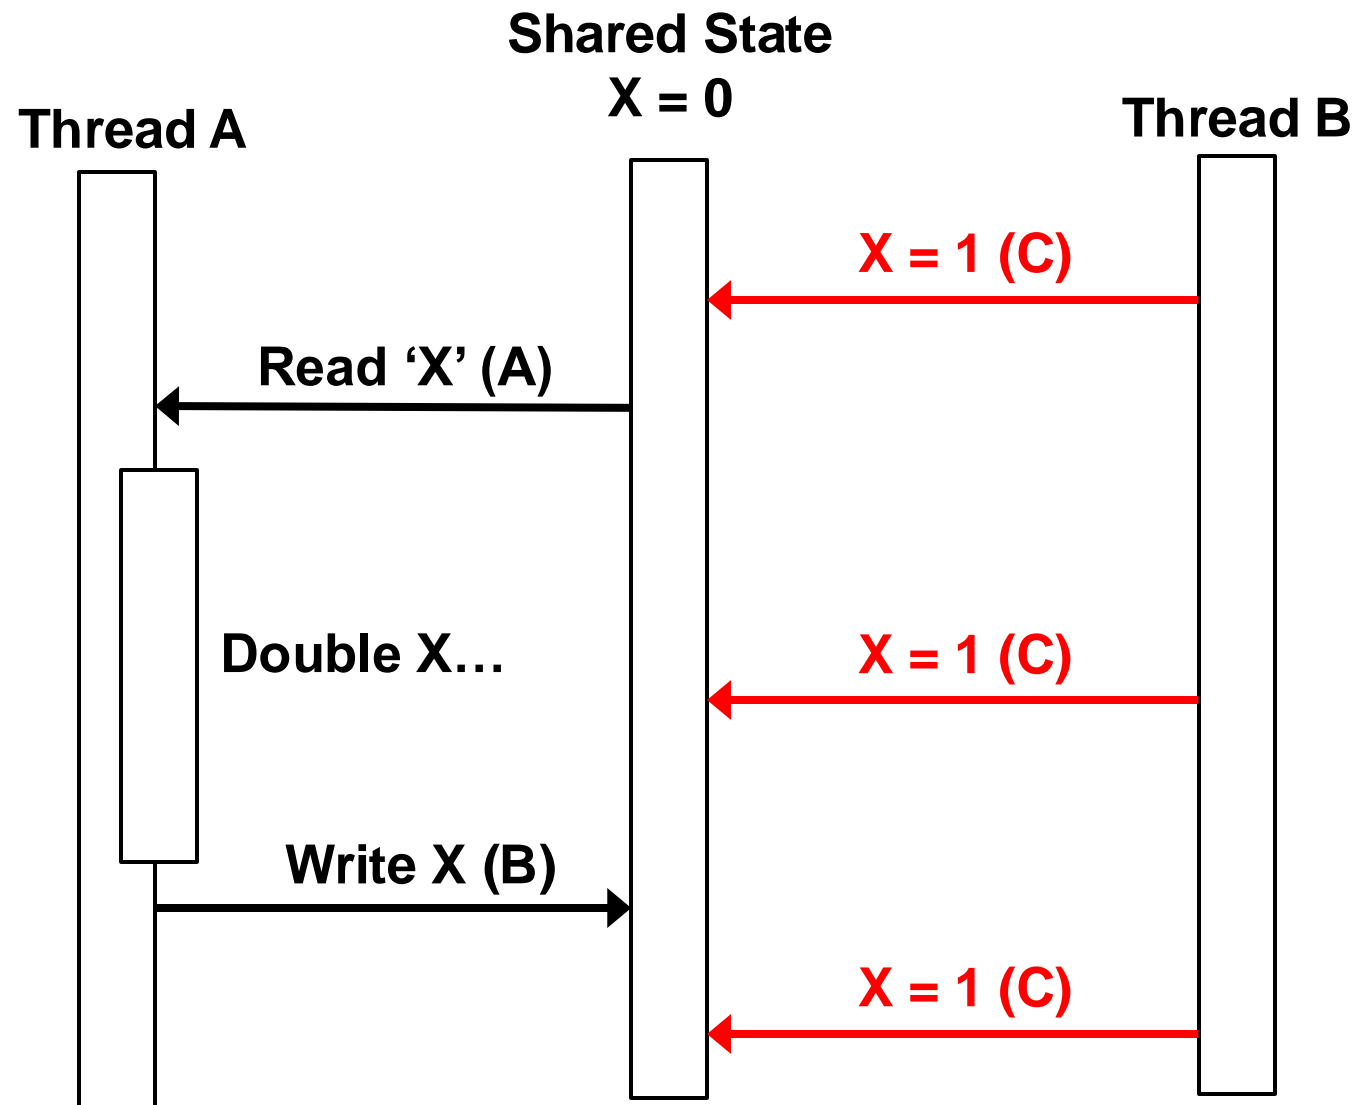
In which order does it happen?
What is the final value of X in each case?

$A \Rightarrow C \Rightarrow B$
 $X = 0$



In which order does it happen?
What is the final value of X in each case?

$A \Rightarrow B \Rightarrow C$
 $X = 1$



In which order does it happen?
What is the final value of X in each case?

*Any of the them?
X = (could be 0, 1 or 2)*

```
class SharedState {
    int x = 0;

    public void doubleX() {
        int x = this.x;
        this.x = x * 2;
    }

    public void setX1() {
        this.x = 1;
    }
}

private void run() throws InterruptedException {
    SharedState someSharedState = new SharedState();
    Thread threadA = new Thread(() -> {
        someSharedState.doubleX();
    });
    Thread threadB = new Thread(() -> {
        someSharedState.setX1();
    });

    threadA.start(); threadB.start();
    threadA.join(); threadB.join();
    System.out.print(String.format("X = %d . ", someSharedState.x));
}

public static void main(String args[]) throws InterruptedException {
    for (int i = 0; i < 100; i++) {
        new BrokenExample().run();
    }
}
```

```

class SharedState {
    int x = 0;

    public void doubleX() {
        int x = this.x;
        this.x = x * 2;
    }

    public void setX1() {
        this.x = 1;
    }
}

```

Output:

```

X = 2 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 0 . X = 1 . X = 1 .
X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 .
X = 1 . X = 1 . X = 2 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 .
X = 1 . X = 1 . X = 1 . X = 0 . X = 1 . X = 2 . X = 1 . X = 1 . X = 2 . X = 1 .
X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 0 . X = 1 . X = 1 .
X = 1 . X = 0 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 .
X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 .
X = 1 . X = 2 . X = 1 . X = 1 . X = 1 . X = 0 . X = 1 . X = 1 . X = 1 . X = 1 .
X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 .
X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 .

```

```

for (int i = 0; i < 100; i++) {
    new BrokenExample().run();
}
}

```

```

class SharedState {
    int x = 0;

    public void doubleX() {
        int x = this.x;
        this.x = x * 2;
    }

    public void setX1() {
        this.x = 1;
    }
}

```

Output:

X = 2 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 0 . X = 1 . X = 1 .
 X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 .
 X = 1 . X = 1 . X = 2 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 .
 X = 1 . X = 1 . X = 1 . X = 0 . X = 1 . X = 2 . X = 1 . X = 1 . X = 2 . X = 1 .
 X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 0 . X = 1 . X = 1 .
 X = 1 . X = 0 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 .
 X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 .
 X = 1 . X = 2 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 .
 X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 .
 X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 . X = 1 .

Program Behaviour is undefined!

**We need to control the ordering of
(i.e. *synchronize*) the events...**


```

for (int i = 0; i < 100; i++) {
    new BrokenExample().run();
}
}


```

Single-Track Rails – for max 1 train...


Train 1




Train 2



Depending on the speed of the
trains, when they set off etc. etc....



Train 1




**If the trains (threads) enter the
“critical section” of track (code) at the
same time – an accident can happen.**

How can we prevent this?

Critical Section



Train 2






See: [https://en.wikipedia.org/wiki/Token_\(railway_signalling\)](https://en.wikipedia.org/wiki/Token_(railway_signalling))



**In Programming:
“Monitor” or “Lock”**


Train 1





Critical Section




Train 2









Train 1




Train 2 (waiting)




Train 1




Train 2




Train 1




Train 2



Train 1



Train 2



The synchronized keyword

```
class SharedState {  
    int x = 0;  
  
    public synchronized void doubleX() {  
        int x = this.x;  
        this.x = x * 2;  
    }  
  
    public synchronized void setX1(){  
        this.x = 1;  
    }  
}
```

We can add the **synchronized** keyword to methods, as shown.

=> A thread can only proceed with the method when it holds the monitor or lock ().

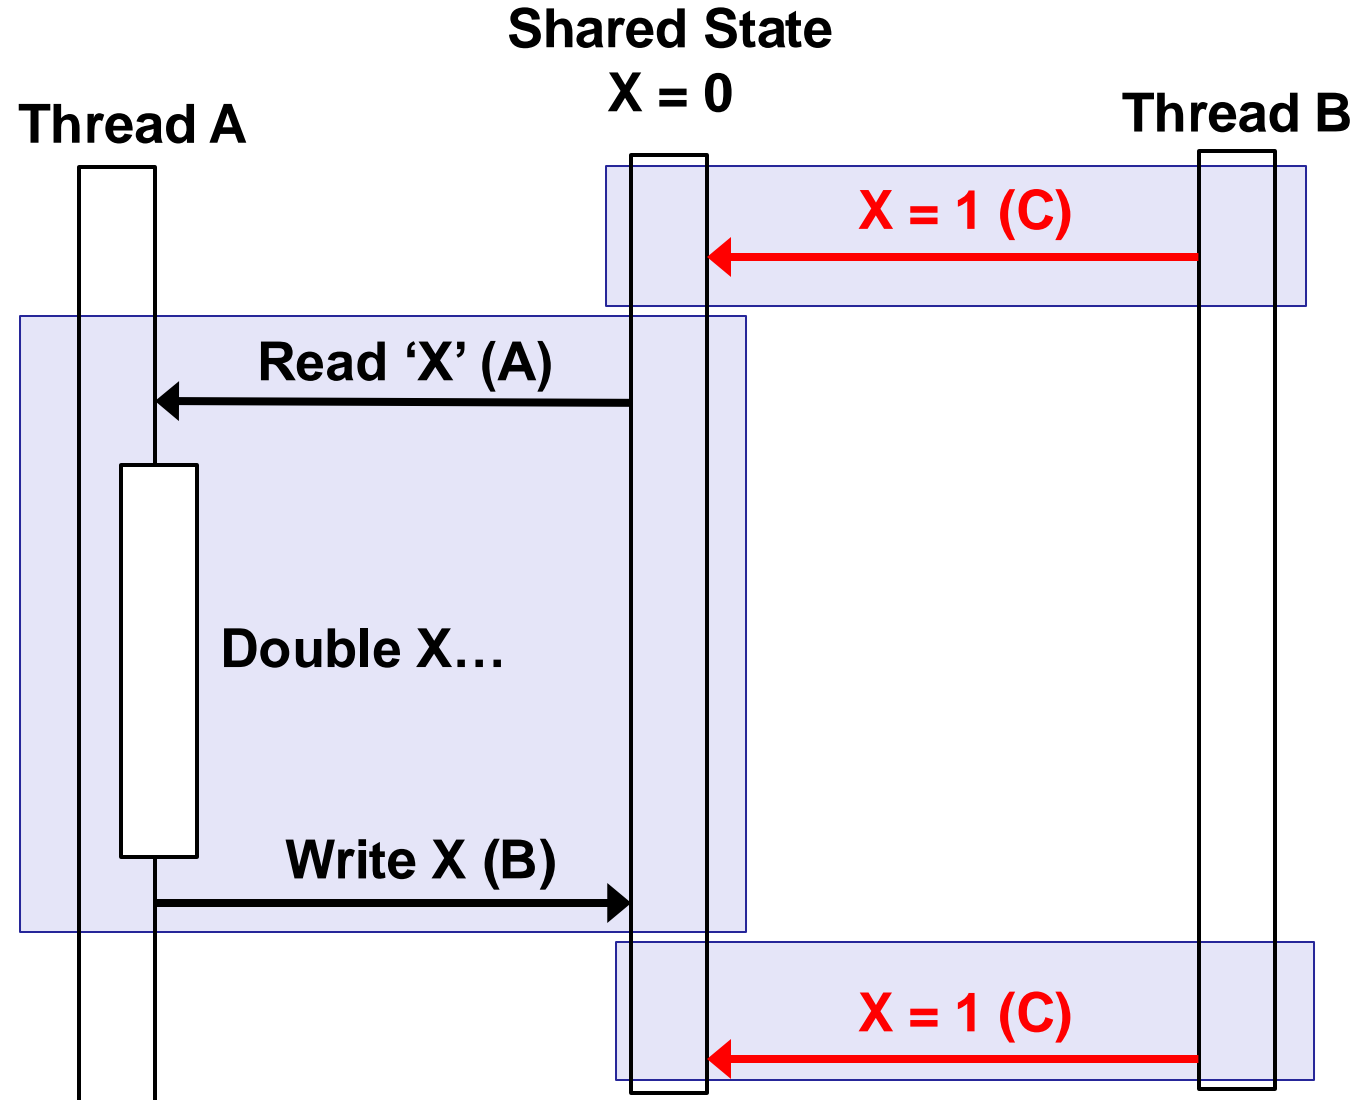
“it is not possible for two invocations of synchronized methods on the same object to interleave.”

When one thread is executing a synchronized method for an object, all other threads that invoke synchronized methods for the same object block (suspend execution) until the first thread is done with the object.”

Source:

<https://docs.oracle.com/javase/tutorial/essential/concurrency/syncmeth.html>

With the synchronized keyword:



Now in what order can it happen?

We can have:

$(C) \Rightarrow (A \Rightarrow B)$
Result: 2

-or-

$(A \Rightarrow B) \Rightarrow (C)$
Result: 1

(C) cannot occur between (A) and (B)!

But we still have a 'race condition' (its still broken)

The synchronized keyword

- The synchronized keyword controls access to all method innovocations on each **object independently**.
- Other threads can invoke synchronized methods **on other objects of the same class**.
- We can use the synchronized keyword when we are doing a complex operation, such as:
 - reading attributes, performing computation, and writing back a new value.
 - this includes **increment/decrement** operations $i++/i--$ (these are really combined read & write operations!)
 - modifying more than one attribute at the same time.
 - or, e.g., operations on lists of objects stored in attributes.
- It prevents multiple threads from interfering with each other, potentially loosing data, overwriting values, leaving attributes inconsistent with eachother, or causing an exception.
- Sometimes, using synchronized on the necessary methods is enough to ensure **thread safety**.
- But, the synchronized keyword does not give any guarantees the order in which threads can execute synchronized methods – its *först till kvarn!*
 - *In our example, the behaviour is still undefined. It's still broken!*

wait() / notifyAll()

- We do not have control over the order in which different threads can invoke synchronized methods.
- For example, with queues, we cannot remove an item from the queue until at least one item is on the queue.
- Sometimes we need a thread to **wait()** for some **condition** so it can complete execution of the synchronized method.
- We can call **notifyAll()** so that other threads to re-check their condition.

wait() / notifyAll()

Now, the doubleX() and setX1() methods will wait() for their conditions to hold.

We allow the two threads to enter the synchronized blocks, but they **still cannot execute concurrently**.

```
class SharedState {
    int x = 0;

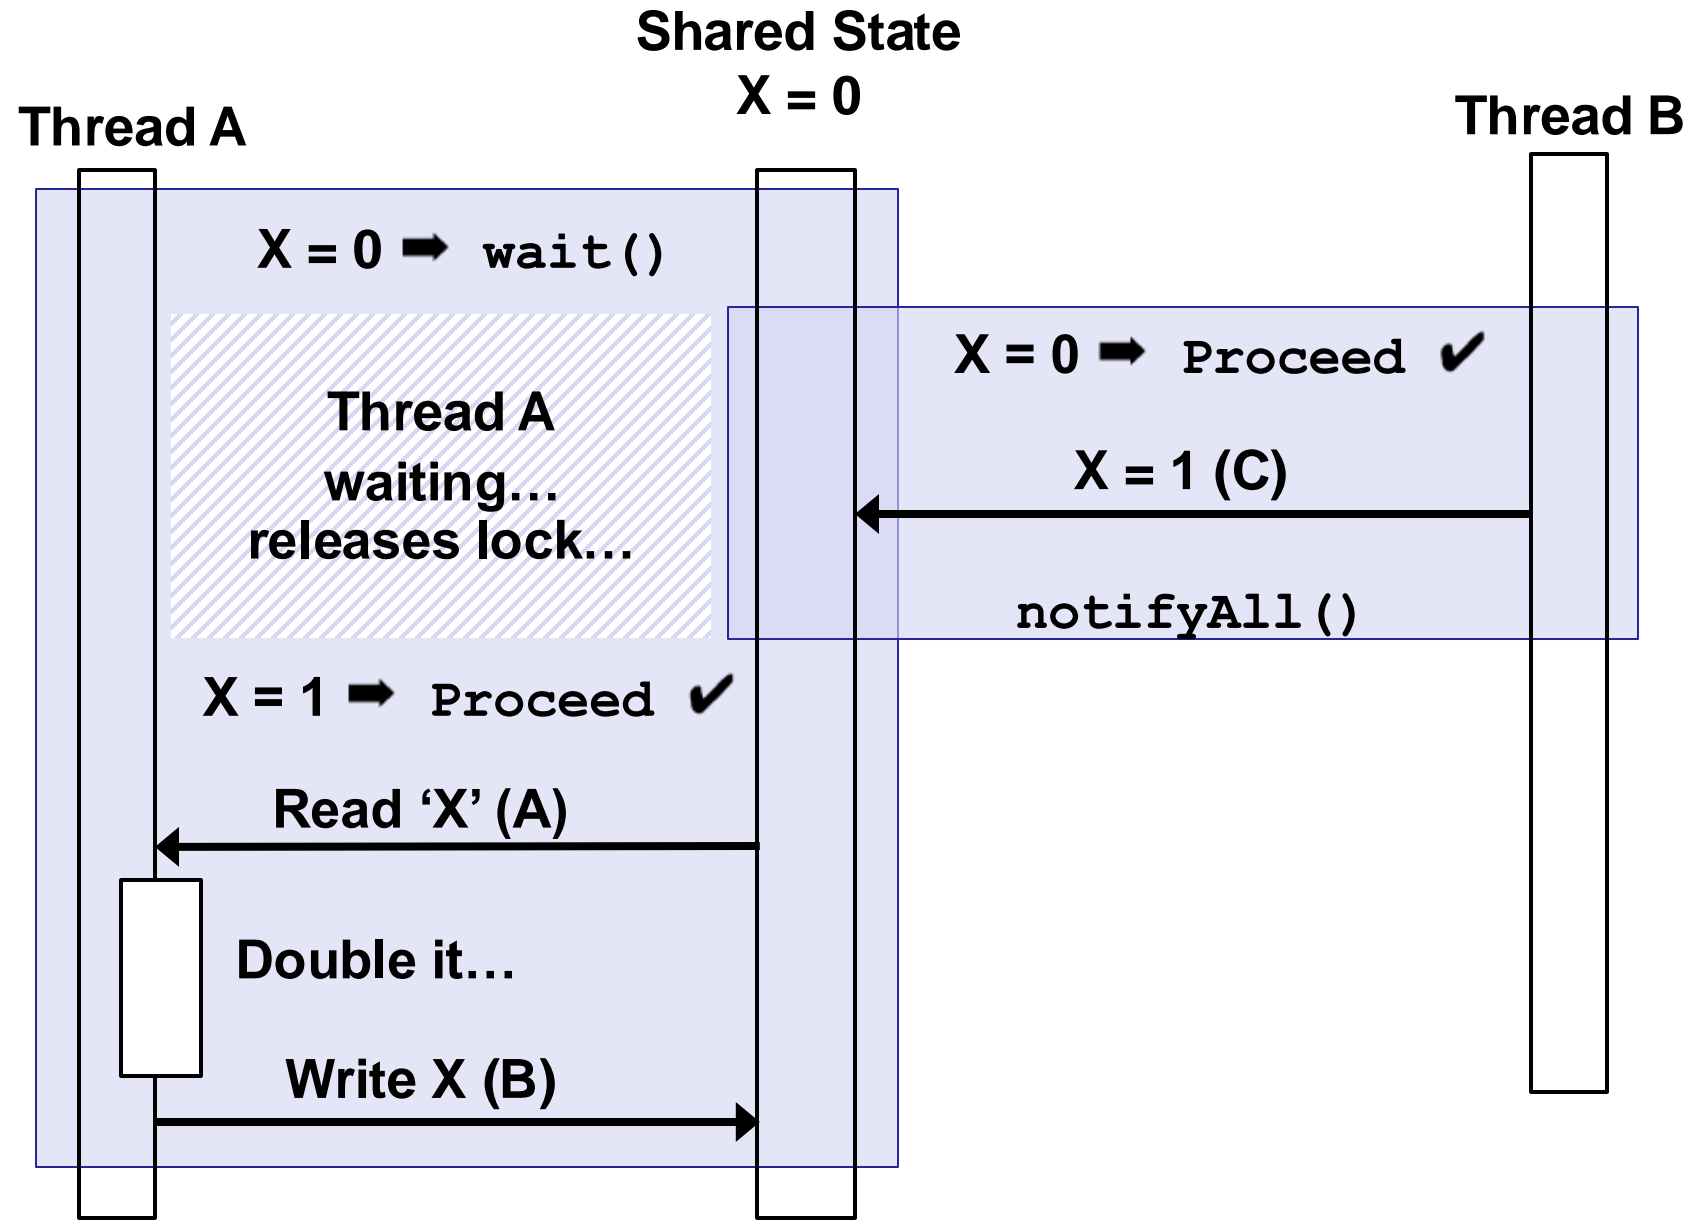
    public synchronized void doubleX() throws InterruptedException {
        // We only want to continue with this method if x == 1.
        while (x != 1) {
            // allow another thread to proceed, and wait for them to call notifyAll()
            this.wait();
        }

        // Now x == 1, we can proceed.
        int x = this.x;
        this.x = x * 2;
        this.notifyAll(); // 'wake up' the other thread when we exit the sync block.
    }

    public synchronized void setX1() throws InterruptedException {
        // We only want to continue with this method if x == 0.
        while (x != 0) {
            // allow another thread to proceed, and wait for them to call notifyAll()
            this.wait();
        }

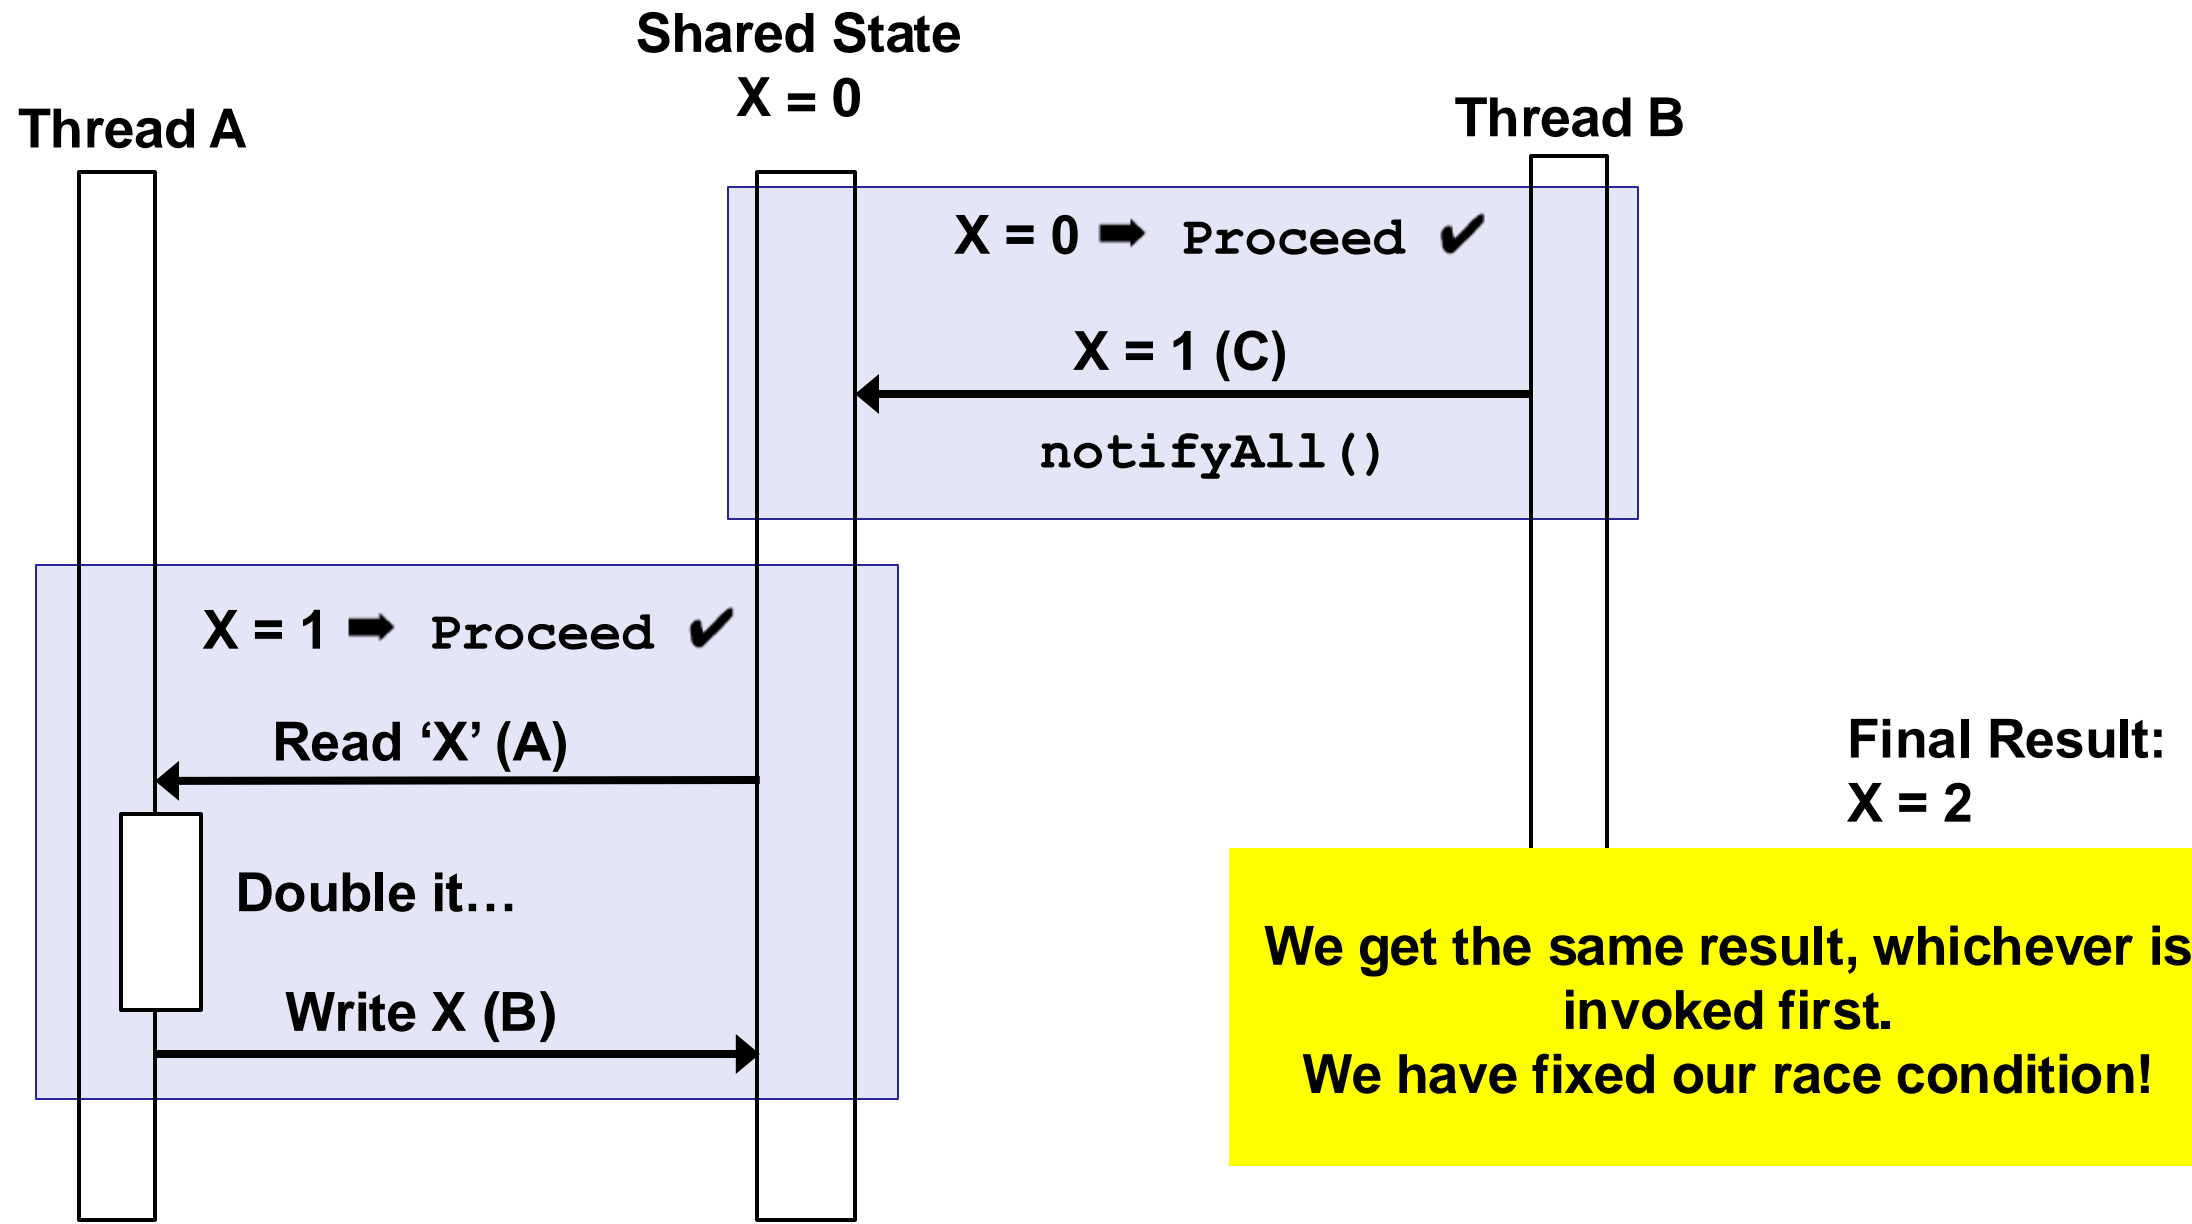
        // Now x == 0, we can proceed.
        this.x = 1;
        this.notifyAll(); // 'wake up' the other thread when we exit the sync block.
    }
}
```

DoubleX() invoked first...



Final Result:
 $X = 2$

SetX1() invoked first...



wait() / notifyAll()

- wait() and notifyAll() can only be used in **synchronized blocks** (otherwise we get an exception).
- wait() and notifyAll() are implemented on the **Object class**.
- We should always **check the condition again after the wait()** – we should not assume the condition now holds just because of the notifyAll().
 - If not, we simply wait() again and check the next time.
 - For this reason, we should always use wait() **inside a loop**.
 - This is an example of a **design pattern**.
- The **notifyAll()** actually only happens when the thread exits the synchronized block (i.e. the method returns).
 - By convention, we put it at the end.
- **notifyAll()** notifies all waiting threads. Still only one thread at a time!
- There is also **notify ()** we can use, in some advanced scenarios this might give better performance.

Design Pattern for wait()/notifyAll()

```
public synchronized void someMethod() throws InterruptedException {  
  
    while (<<condition is false>>) {  
        // release lock, allow another thread to proceed  
        // and wait for them to call notifyAll()  
        this.wait();  
    }  
  
    // Now, <<condition is true>>, and we have the lock.  
    // We can safely make our changes to the state ...  
  
    this.notifyAll(); // 'wake up' the other thread when we exit the method.  
}
```

Synchronized Blocks

- Sometimes, we don't need the whole method to be a synchronized, just part of it.
- We can create a **synchronized block** within a method:

```
synchronized(this) {  
    ....  
}
```

- We can use any object as the lock!
 - Typically we use 'this'.
- The same as putting synchronized on the method.
 - Which implicitly uses 'this' as the lock.
 - We can still use notifyAll() / wait()



Break?

Classes in libraries are often not synchronized (i.e. not thread safe)

java.util.HashMap<K,V>

Note that this implementation is not synchronized. If multiple threads access a hash map concurrently, and at least one of the threads modifies the map structurally, *it must be synchronized externally*. (A structural modification is any operation that adds or deletes one or more mappings; merely changing the value associated with a key that an instance already contains is not a structural modification.)

Source: <https://docs.oracle.com/javase/8/docs/api/java/util/HashMap.html>

Synchronization is expensive – we only use it when necessary.

One Solution: Facade Design Pattern

To synchronize such a class, you can:

1. Set an instance of the class to be private attribute of a new class
2. write synchronized methods that use the instance.


```
public class SynchronizedHashMap<K,V> {  
    private HashMap<K,V> map = new HashMap<K,V>();  
  
    public synchronized V put(K key, V value) {  
        return map.put(key, value);  
    }  
  
    public synchronized V remove(K key) {  
        return map.remove(key);  
    }  
  
    public synchronized void clear() {  
        map.clear();  
    }  
  
    public synchronized V get(K key) {  
        return map.get(key);  
    }  
}
```

Thread Termination

- To terminate the application, **all the threads** need to terminate.
- We want to do this **cleanly** so that all important tasks finish before the application exits – for example saving the user’s work.
- We can wait for a thread to finish by calling **join()** from another thread.
 - This call ‘blocks’ until the other thread is terminated.
- We can also call **interrupt()** on the thread object.
- See: <https://docs.oracle.com/javase/8/docs/api/java/lang/Thread.html>
- A common design pattern for worker threads is to use a **loop**:
 - A queue of incoming tasks or items to process.
 - If the queue is empty, we **wait()** for more tasks.
 - Also check the status of the **interrupted()** flag, to break the loop and terminate the thread.
 - Can also have a boolean “closing” flag, to break the loop and terminate the thread.

Queue (kö) Design Pattern

- A **queue (kö)** is a familiar concept.
- A queue is usually **FIFO** (“first-in-first-out”)
- **LIFO** (“last in first out”), we call a **stack**.



- Components which **add** (or **put**) items to the queue are sometimes called a **producers**
- Components which **remove** (or **get, pop**) items from the queue are sometimes called a **consumers**
- There might be **more than one** producer/consumer.
- We often use queues to send data between threads. This is a form of **buffer**.
- This pattern allows the producer thread(s) to continue working while consumer thread(s) process items.

Buffer<T>

```
public class Buffer<T> {  
  
    // LinkedList<T> is not synchronized (thread safe)  
    private final LinkedList<T> buffer = new LinkedList<T>();  
  
    public synchronized void put(T o) {  
        buffer.addLast(o);  
        notifyAll();  
    }  
  
    public synchronized T get() throws InterruptedException {  
        while(buffer.isEmpty()) {  
            wait();  
        }  
        return buffer.removeFirst();  
    }  
  
    public int size() {  
        return buffer.size();  
    }  
}
```

Buffer<T> Example

See:

Buffer.java

Consumer.java

Producer.java

BufferDemo.java

Exercise: adapt the example there is a queue of *Runnable* objects, and have the consumers execute them.

Can you adapt your solution to use callbacks with the results?

Writing Multi-Threaded Code

- With **single-threaded code**, if it works (and we have tests) – we feel like we're done!
- With **multi-threaded code**, just because it works, it might still be broken.
 - Even if we run our program several times, and we get the correct output, this does not necessarily mean we have thread-safe code.
 - Running the same code later on, or on a different machine, may produce different results.
 - Maybe there will only be a problem .01% of the time!
 - It could be that future modifications to the code will expose existing problems with thread safety, which were just avoided with good luck.
- Nor can we easily write unit tests for thread safety.
- There are tools for code analysis (outside the scope of the course).
- We imagine how different threads can interact with our objects – and call methods in all possible different orders.
- Use the debugger!
- Print out the name of the current thread. (can also **assert** which thread we are on)

Questions?